

Evented I/O based web servers, explained using bunnies

Simon Willison

Part of a talk given at Full Frontal 2009



**Your Web Server
(using a bunny)**



Your Web Server
(using a bunny)

Single threaded (one bunny), so can only
handle one request at a time

Happy hamster



(The hamsters are using web browsers to visit your site)



Your Web Server
(using a bunny)

Impatient hamsters

5 bunnies = can handle 5 requests at
the same time



Happy hamsters



Your Web Server
(using threads,
aka bunnies)



Long running operations cause a thread to block, causing requests to build up in a queue



fetching a web API
(2 seconds)



Your Web Server
(using threads,
aka bunnies)



Impatient hamsters



fetching a web API
(2 seconds)



uploading an image
(3 seconds)



fetching a web API
(2 seconds)




fetching a web API
(2 seconds)



comet long polling
(10 seconds)

Replace the bunnies with a single hyperactive squid. The squid runs up and down as fast as it can dealing with each hamster in turn. It fires off any long running I/O operations and then moves on to the next hamster. When the I/O operation reports progress, it does a little more work on behalf of the corresponding hamster.



Your Web Server
(event loop, aka
hyperactive squid)

Bad code

- `rows = database.fetch(category = 'news')`
- `template = read_file('homepage.html')`
- `json = fetch_url('http://.../')`

These functions block and wait for results - blocking the squid and causing the entire event loop (and hence server) to pause until they complete

Good code

- `database.fetch(category = 'news', callback)`
- `read_file('homepage.html', callback)`
- `fetch_url('http://.../', callback)`

These functions specify a callback to be executed as soon as the I/O operation completes